


Practical Password Cracking

The First 500 Million

jamie.riden@ioactive.com

IOActive[®]





Acknowledgments

- None of this would be possible without
 - Hashcat itself
 - John the Ripper
 - Impacket
 - Wordlists (Seclists, probabilistic passwords, Crackstation, hashes.org, etc.)
 - D3.js and bokeh (graphing)



- Built/improved multi-user cracking systems at two UK pentest companies
- Cracking other people's passwords since 2003, red and blue "teams".
- Dev, sysadmin, blue team, now pentester
- From embedded systems up to a Cray
- Erdős number of 4



Why Bother?

- Naïve approach – rockyou.txt & nsav2dive.rule
- Actually, not bad, but we can do better
 - 40+ common hash types to remember
 - Format conversions
 - Pen tests have hard time limits
- “Improving the Security of Your Site by Breaking Into it”
- “31.3. To maximise test coverage the following types of automatic test equipment shall be used where relevant: password cracking tools; “ – CHECK guidance



Example Hashes

```
6db61eaa9962d0ef0fd6e7676c9f8e68      [NTLM or MD5]
{SHA}AEkejO/OIdcoo8beMltXFfsSUoaE=      [Netscape LDAP SHA1]
25f7edf2e81c5c5ae89c86d8d0581ed3f1d0d44657078204a551717f61d52f23
df1bxxx8d32ba0edacf047ca0349d05cf      [SHA512, sometimes SHA2-512]
{SSHA}jmu2X5k8vOJa5XZGeUa9TlrQlpEw     [Netscape LDAP salted SHA1]
1.ke4gjHbX/Xg                            [DEScrypt, 8 char max]
$1$LAr1jfuu$ztACi8wvC1X/cMk7B8Owg/      [md5crypt]
$6$0rqfz1DPSQDWk$gOk9W85/qwUYiDs/hNsgux.Z1GOQVRHnHFNY1dOWPKIzYxk
gLMZcbGoU37xt06xiMzM2.RaPYh/e8niiWtDAQ. [sha512crypt]
$2a$10$XFX4Pyb/KjTgaxH4cFHhbOsP8PsECxxxRw7zwRWTYJi6 [bcrypt]
$9$DDHX/GcKrS.OUZ$OqUtJrbh2eOaNZGWBCxxxD.9r.gpVD2os [scrypt]
```



- Bad Hashes (fast, unsalted hashes)
- Good Hashes (slow, salted hashes)
- Hashcrack script/tools -
<http://github.com/blacktraffic/hashcrack>
- Offensive uses
- Defensive uses
- Cracking 500 million hashes
- Quick spin against CMIYC 2015 hashes (CTF)



Bad Hashes

- MD5 hash: “secret” -> 5ebe2294ecd0e0f08eab7690d2a6ee69
- MD5 is a quick hash function mapping anything to a 128-bit value.
- Very quick – 35 billion guesses / sec.
- Can build lookup table (hash, password)
- Ignore other bad properties of MD5
- NTLM is basically widechar MD4. (LM is worse.)



Bad Hashes

- No home-brew crypto please
- Adding “secret” static data to your MD5 doesn’t help as much as you think.
- Improve hash function before you add ‘pepper’. (per app secret)
- Composing bad functions (md5, sha1, etc) doesn’t really help
- Avoid poor Random Number Generators
- **Simpler to use good library function**



- Blowfish, or argon2, or PBKDF2 etc.

```
echo password_hash("test", PASSWORD_DEFAULT);  
$2y$10$.vGA1O9wmRjrwAVXD98HNOgsNpDczlqm3Jq7KnE  
d1rVAGv3Fykk1a
```

- Salt, cost parameter, very slow – 20 thousand guesses / sec with cost 5.
- Random salt prevents lookup table creation




- Guesses hash type (regular expressions mappings)
- Tries a decode if you need one (mostly via JtR scripts)
- Tries some reasonable dictionaries, rules
- Configurable – mapping files are data
- Can override for specific requirements
- AGPL



Hashcrack Script - Motivation

- “Laziness, hubris, impatience”
- Codify some attack techniques
- Measure, improve
- “front loading”
- Don’t make me think
- Often stuck in machine rooms

Cumulative cracked ratio over 2018





Hashcrack Script – Constraints & Goals

- 5 main attack types with hashcat:
 - Hybrid (dict+rules), cross-product, masks (3)
- Keep the GPU busy:
 - Fast hashes need more than just dictionary
- Efficiency
 - try not to duplicate guesses (dict + rules)
 - Go for greatest speed/density first
- Re-use pwdump LM cracks for NTLM
 - yes, we still see LM! LM password is uppercase NTLM password



Sketch of an Attack Workflow

Fast hash like pwdump (uid:LM:NTLM)

```
python3 hashcrack.py -i pwdump.txt
```

1. Crack LM first using incremental up to 7 chars
2. Use this as crib for NTLM (all case permutations)
3. Incremental up to 8 for NTLM
4. Run found.txt with various options, if present
 1. -a6 -i ?a?a (found.txt with all 1,2 char suffixes)
 2. -a1 last3-4.txt (with common 3-4 char suffixes)
5. Decent sized dict with l33tpasspro rules



Sketch of an Attack Workflow.

Slow Hash like bcrypt:

```
python3 hashcrack.py -i tests/bcrypt.txt
```

0. (manual) get trunk hashcat from github & compile – 40% quicker
 1. Small crib file with company name etc.
 2. Previously cracked passwords with no rules
 3. Medium dictionary, no rules



Steal hashes via


- Responder
- Kerberoasting
- Internal Monologue
- Physical theft of device/disk/vmdk etc.
- smb:// URIs and similar leaks
- Dumping NTDS or local registry hives
- Cisco/Juniper/etc configs
- Office docs, zip files, Bitlocker volumes ...



- Once you have the hash, cracking is opaque to the target
- Generally lots of things to try
- Never know what one success will achieve
- Get something started quickly while you have a proper look around
- Once you get some passwords cracked, try them on other things (“crib”)



Guess Speed (Ideal)



* LM is awful for other reasons too



Red Team – Use Cases

```
python3 hashcrack.py -i ifm.zip
```

[invokes impacket/secretsdump.py to decode]

```
python3 hashcrack.py -i Responder.db
```

[uses sqlite3 library to query the db file]

```
python3 hashcrack.py -i salaries.xlsx
```

[runs jtr script office2john.py]



Red Team – Use Cases


```
$ python3 hashcrack.py -i tests/test-abc.docx -d hashcat.txt
Reading file: /root/hashcrack/tests/test-abc.docx
RUN: python2 office2john.py tests/test-abc.docx > tests/test-
abc.docx.tmp2
RUN: ./hashcat64.bin -a0 -m 9600 tests/test-abc.docx.tmp2
hashcat.txt -r rules/l33t64.rule --loopback -O -w4
hashcat (v5.1.0) starting...
```

```
$office$*2013*100000*256*16*xxxa7f38d98d12:abc
```



“Boring” Audit Stuff – Pentest or Blue Team

- Export domain passwords
- Crack the hashes
- Uses LM cracks with case permutations, if present
- Draw some graphs if you like
- Update policy, make users change password





Handling painful stuff – e.g. Oracle

```
$ python hashcrack.py -i oracle-hashes.txt -t oracle --  
show
```

```
CTXSYS:S:54C7Cxxx;T:A10609DFAE71F2xxx1C27D2E4E69CB6BAB96  
EDE06F795842CE7
```

Found 112

```
54c7ccf5def04820dxxx7d61a6b:e5ec99bc210308f58747:CTXSYS
```

Found 12300

```
A10609DFAE71xxx9CB6BAB96EDE06F795842CE7:CTXSYS
```





Blue Team – Quick Wins

- Stop users from choosing
 - Really weak passwords
 - Widely compromised passwords
 - Passwords based on dictionary words
- Rate-limit password guesses somehow.
- Back off to CAPTCHA for potential brute-force or credential stuffing attacks
- Crack your own password hashes



Stats from the 500 million cracked





Querying Troy Hunt's Compromised List

```
$ curl  
"https://haveibeenpwned.com/api/v2/pwn  
edpassword/5baa61e4c9b93f3f0682250b6cf  
8331b7ee68fd8
```

```
HTTP/2 200 [found in a breach]
```

```
$ curl  
"https://haveibeenpwned.com/api/v2/pwn  
edpassword/d5c0a7e48a8047f6c59abba1db3  
b365ecdaf3663
```

```
HTTP/2 404 [not found in breach]
```





What was cracked and when

Can derive from
hashcat status
output

Minimise the
area under this
curve:

kill 'weak' ones
(lhs)



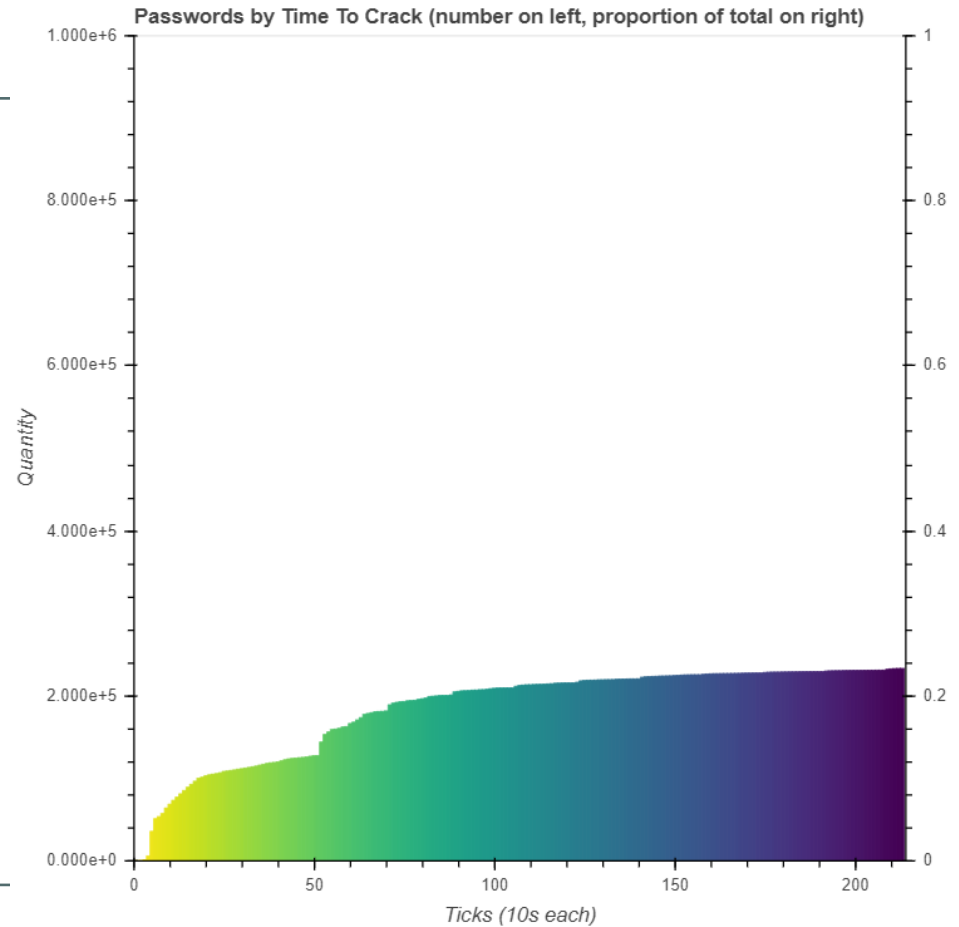


What was cracked and when

First million from HIBP


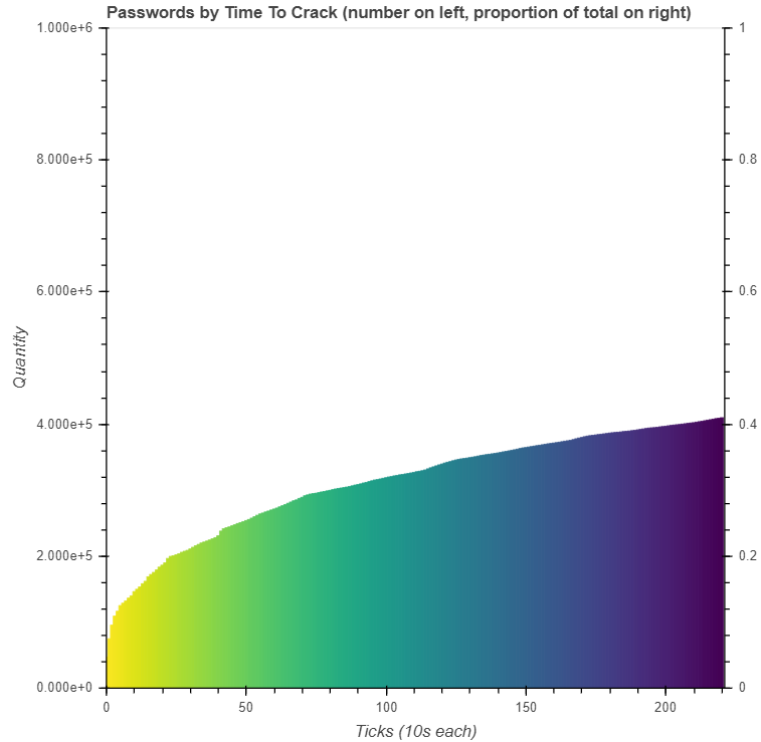
Cracking on laptop

Incremental first then
hybrid dict and rules





What was cracked and when – different rules, effect of pot file



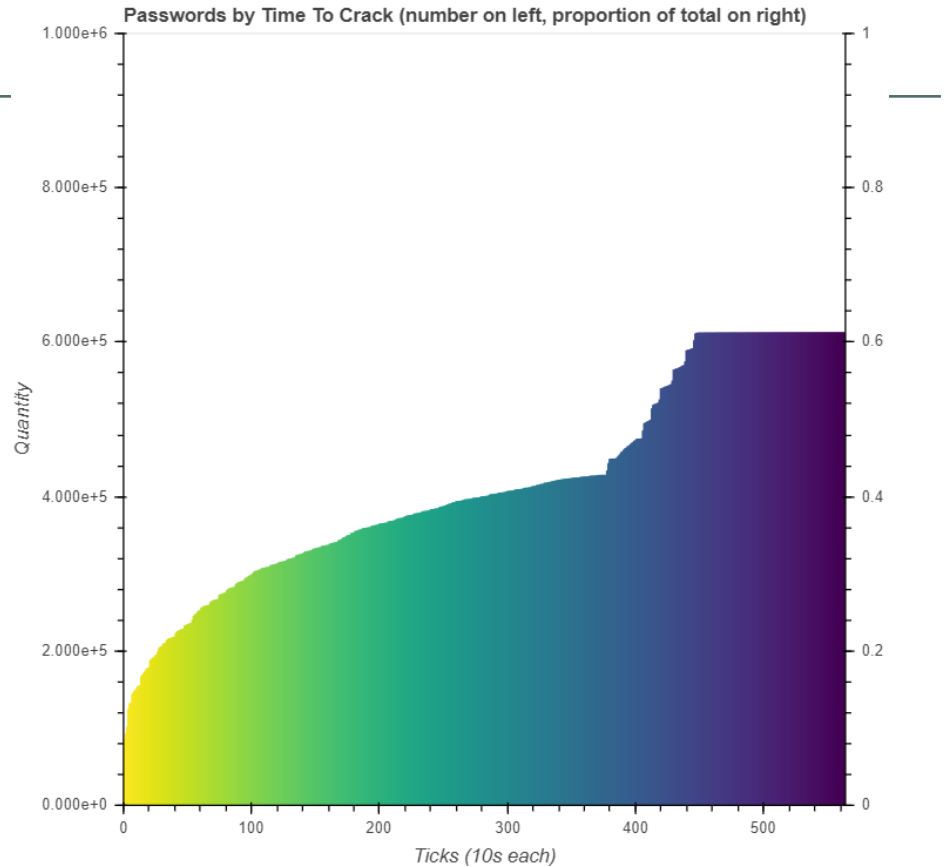


What was cracked and when

Now, on a 1080 Ti

1st: Top 258 Million

2nd: incremental up to 9
- starts about tick 400





What was cracked and when – How to Generate

`$ python3 hashcrack.py -i testcase.txt --status greyarea2`
(writes the status updates to a file 'greyarea2')

`$ python graph-by-quality.py greyarea2`
(uses bokeh to plot a graph from the status file – probably will be in the hashcat-5.1.0/ subdirectory)

--status is the hashcat flag used to generate this every 10s



Design of a Password Cracking System

power, cooling, contention, latency, throughput

Entry level – laptop. **Do not overheat.**

Modest tower with 1080 Ti (greyarea)

Homebrew multi-card setup

“COTS” hw - Brutalis (<https://terahash.com/>)

Clustered system, e.g. hashtopolis

Multi-user cluster, custom code





Design of a Password Cracking System

Bulk of this was:

- 1080 Ti (grey area)
- Dell 7510 laptop

Don't skimp on the card

1080 Ti is 250W, cooling is important





Cracking 500 Million Passwords – Getting The Data

<https://haveibeenpwned.com/Passwords>

```
wget https://downloads.pwnedpasswords.com/passwords/pwned-passwords-ntlm-ordered-by-hash-v4.7z
```

```
7z x pwned-passwords-ntlm-ordered-by-hash-v4.7z
```

```
cat pwned-passwords-ntlm-ordered-by-hash-v4.txt | cut -f 1 -d':' > pwned.txt
```

Tweak the hashcrack.py code with --bitmap-max=26



Cracking 500 Million Passwords – Split / Crack / Merge

```
$ split -l 50000000 -d pwned.txt      # 50 mil at a go
```

```
#!/bin/bash
```

```
python3 hashcrack.py -i x00 -t ntlm --pot troy1.pot --remove >/dev/null
```

```
... repeat ...
```

```
cat x?? > pwned-remaining1.txt
```

```
rm -f x??
```

Need >/dev/null for the first couple of phases, too much output



Attack Types

-a0 dict + rules. Rules often needed for throughput
-a0 Top258Million.txt -r rules/l33tnsa.rule

-a3 incremental with given mask, e.g. ?u?l?l?l?d?s

or with hcmask file: ?u?l?d,?l?d,?l?d?s,?1?2?2?2?2?3

-a6, -a7 left and right mask -a6 found.txt ?a?a?a -i

-a1 leftdict rightdict *aka "information supercollider"*



Attack Types

-a1

Left dict:	your	my	the
Right dict:	cat	dog	moose
	your	my	the
cat	yourcat	mycat	thecat
dog	yourdog	mydog	thedog
moose	yourmoose	mymoose	themoose



Attack Types and Flags

`--loopback -O -w4` # optimise for Linux (-w3 for Windows), feedback

`-a1` allows cross-product. combinator and combinator 3 do similar. Or write your own script:

```
$ python leet.py words.txt | ./hashcat64.bin -m 1000  
/root/hashcrack/pwned-remaining3.txt -a0 -O -w4 -r  
rules/best64.rule --bitmap-max=26 -r rules/best64.rule
```

OMEN & PRINCE & PCFG as candidate generators as well



STDIN, use with generator script, e.g. combinator or

```
$ python scripts/leetify.py dict/Top32Million-  
probable.txt | ./hashcat-5.1.0/hashcat64.bin -m 1000 -  
O -w4 pro-hashes.1.txt -r rules/best64.rule
```



Attack Types - OMEN

#build model

```
./createNG --iPwdList ../dict/cracked-passwords.txt
```

#invoke model to generate candidates

```
./enumNG -p | ../hashcat-5.1.0/hashcat64.bin -m 1000 -O -w4 street-hashes.1.txt -r  
../rules/nsav2dive.rule
```

beb68df17dfbe162212ff8a5d8a254da:Pickey21

31e25e76883ea48224473670cd84aa91:Aller13*

a78be982ee90685cbf8e3b31bd14b4db:Sandand09



Attack Types - Combinator

```
./hashcat-utils-1.9/bin/combinator3.bin top95k.txt  
top95k.txt top95k.txt | ./hashcat-5.1.0/hashcat64.bin -  
m 1000 -O -w4 /root/bits/remaining5.txt -r  
rules/nsav2dive.rule
```

```
1d227xxc91ed5e4fb073:1357iloveyou1357
```



With dict/words.txt three times :

- f1d99b1c353d01xx: blueberrycarrot
- e25d103708a7eaxxc: blueberryfudge
- 99ed45f86ffe309xx: bluebirdpelican123
- 6af87dc65d6cxx6517f: bluefisheggs



With dict/top1k.txt three times :

- onrunner12345
- 43211234asd
- sunshinelovepark



Beyond 7-bit ASCII

LM seems to depend on codepage.

NTLM is basically widechar MD4. (e.g. -m 900 is raw mode)

See: <https://www.blackhillsinfosec.com/cracking-passwords-with-umlauts/>

SHA1 and related UNIX hashes like SHA512crypt are generally UTF-8 – web sites may choose their own encoding.



Attack Types

`hashcrack.py -i hashfile.txt [-d dict] [-r rules]`

`hashcrack.py -i hashfile.txt --mask hcmaskfile or literal`

`hashcrack.py -i hashfile.txt -d dict -lmask <mask>`

`hashcrack.py -i hashfile.txt -d dict -rmask <mask>`

`hashcrack.py -i hashfile.txt -d leftdict -e rightdict`



Cracking 500 Million Passwords – Phases

```
248945621 troy1.pot [defaults -a0]
 96500674 troy2.pot [def incr -a3]
141627978 troy3.pot [breachcompilation]
 2777509 troy4.pot [hcmask incr -a3]
 9121343 troy5.pot [Top 2 Billion]
1090010 troy6.pot [258Mil/last 4, a1]
1557763 troy7.pot [other dict / last]
24902724 troy8.pot [hashes.org]
  68403 troy9.pot [misc ideas]
```



Cracking 500 Million Passwords – Lessons ?

Not a huge amount?

While very much a real world dataset, it's not one you will need to attack in real life.

Do NOT use it as a dictionary for NTLM cracking as is.

First remove anything really weak, saves repetition

```
hashcat64.exe -m 99999 example.txt -O -w3 -a3 --remove -  
i --increment-max=7
```



<https://contest-2015.korelogic.com/downloads.html>

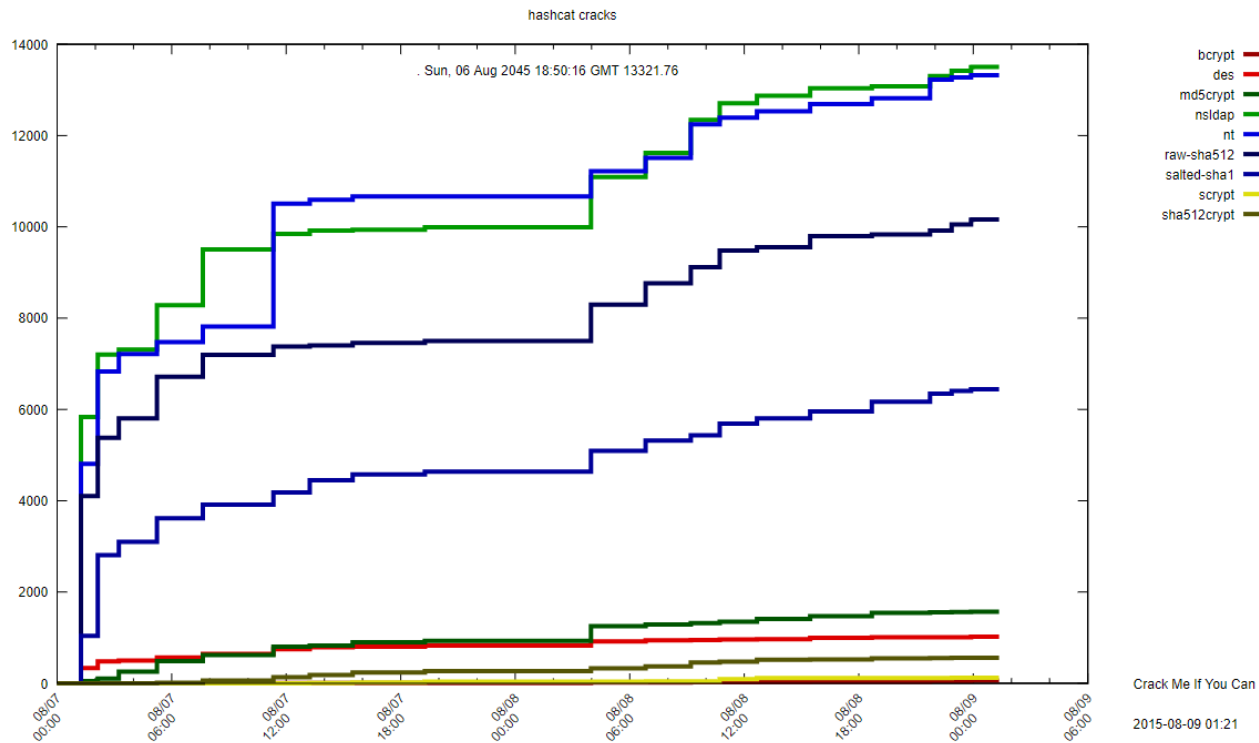
Download, decrypt, unpack

```
$ for i in `ls street` ; do python hashcrack.py -i street/$i ;  
done # laptop, Quadro 500
```

```
$ for i in `ls pro` ; do python hashcrack.py -i pro/$i ; done #  
1080Ti
```



Crack Me If You Can 2015 – Winners





Crack Me If You Can 2015 – Our “phone it in” attempt

After ~48 hours on a single 1080Ti, cancelling some of the very long runs (minimal user intervention).

```
$ for i in `ls cmiyc/pro/*` ; do  
python3 hashcrack.py -i $i -remove -  
-pot cmiyc/pro.pot ; done
```

```
$ wc -l /root/cmiyc/pro.pot  
6252 /root/cmiyc/pro.pot
```




Future Work

More graphs and metrics – this is a root word and rules leading to cracks (edge nodes)

~~Multiple back-ends (JtR~~
~~— for e.g. bitlocker) done~~

PRINCE, OMEN, PCFG etc





References

<https://uncommoncriteria.org/ppc.html> : Supplementary materials.

<https://www.netmux.com/blog/hash-crack-v3> – Like the ‘RTFM book’ but for password cracking. Highly recommended, name clash is accidental

<https://cyberwar.nl/d/1993-FarmerVenema-comp.security.unix-Improving-the-Security-of-Your-Site-by-Breaking-Into-It.pdf> Farmer & Venema 1993

<https://haveibeenpwned.com/Passwords> Troy Hunt’s list as SHA1/NTLM

<https://blog.cynosureprime.com/2017/08/320-million-hashes-exposed.html> Another (and better) attempt at Troy Hunt’s list



References

https://github.com/lakiw/pcfg_cracker – PCFG candidate generator

<https://hashcat.net/wiki/doku.php?id=princeprocessor> - PRINCE

<https://github.com/RUB-SysSec/OMEN> - OMEN

https://hashcat.net/wiki/doku.php?id=hashcat_utils – combinator, etc.



Rules and Dicts

Recommended: nsav2dive.rule

TopNMillion (probabilistic passwords)

Breachcompilation

Hashes.org “left” lists

Improved? “leetified” versions of nsav2dive, best64 and PasswordsPro rulesets.

New: Insertions.rule : any one or two chars inserted at all possible places in a word.

lastN.txt – all suffixes of N letters taken from breachcompilation or similar.

lastN-M.txt – all suffixes between N and M chars long



Live Demo ? ^_(\ツ)_/

```
Command Prompt - test.bat

OpenCL Platform #1: Intel(R) Corporation
=====
* Device #1: Intel(R) UHD Graphics 620, 4095/13036 MB allocatable, 24MCU
* Device #2: Intel(R) Core(TM) i7-8550U CPU @ 1.80GHz, skipped.

OpenCL Platform #2: NVIDIA Corporation
=====
* Device #3: Quadro P500, 512/2048 MB allocatable, 2MCU

Hashes: 1 digests; 1 unique digests, 1 unique salts
Bitmaps: 16 bits, 65536 entries, 0x0000ffff mask, 262144 bytes, 5/13 rotates
Rules: 98629

Applicable optimizers:
* Optimized-Kernel
* Zero-Byte
* Precompute-Init
* Precompute-Merkle-Demgard
* Meet-In-The-Middle
* Early-Skip
* Not-Iterated
* Appended-Salt
* Single-Hash
* Single-Salt
* Raw-Hash

Minimum password length supported by kernel: 0
Maximum password length supported by kernel: 31
Minimum salt length supported by kernel: 0
Maximum salt length supported by kernel: 51

Watchdog: Temperature abort trigger set to 90c

Dictionary cache hit:
* Filename..: C:\Users\jamie\Desktop\hashcrack\dict\Top95Thousand-probable.txt
* Passwords.: 94983
* Bytes.....: 821551
* Keyspace..: 9368078307

The wordlist or mask that you are using is too small.
This means that hashcat cannot use the full parallel power of your device(s).
Unless you supply more work, your cracking speed will drop.
For tips on supplying more work, see: https://hashcat.net/faq/morework

Approaching final keyspace - workload adjusted.

[s]tatus [p]ause [b]ypass [c]heckpoint [q]uit =>
```



Questions
